

## Problem Statement

The problem statement is to implement a State-based Conflict-free replicated data type(CvRDT) system using the set data structure.

Two endpoints to be made available by each node to receive insert and remove items requests from the main\_set that is stored locally.

- When an item is inserted to the set using the /add endpoint, the item must reflect in other nodes' set too.
- When an item is removed from the set using the /rem endpoint, the item must be removed from the other nodes' set too.

## Proof of Correctness

In the case of implementing a CvRDT, there are few important properties needs to be hold, they are-

### Convergence

This property ensures that eventually all replicas or nodes will end up in the same state. In this case, the main\_set will be the same across all the replicas.

### Proof

This property ensures that eventually all replicas or nodes will end up in the same state. In this case, the main\_set will be the same across all replicas.

- Each replica maintains three ordered set data structure variables: main\_set, add\_set, and rem\_set.
  1. main\_set stores the final state locally. Insertions or deletions requests from users at the /add and /rem endpoints are applied directly to the main\_set.
  2. add\_set stores the items being inserted into the main\_set. Whenever an item is inserted into the main\_set, it's also inserted into the add\_set.
  3. rem\_set stores the items being removed from the main\_set. Whenever an item is removed from the main\_set, it's inserted into the rem\_set.
- The add\_set and rem\_set are fully broadcasted over UDP using another thread every 1 second, as we're implementing a state-based CRDT.
- When a replica receives a UDP packet, it first iterates over the received add\_set and inserts items from it into the main\_set based on two conditions:
  1. The item is not already present in its local add\_set (ensuring that the item is not present in the main\_set).
  2. The item is not present in the received rem\_set over UDP (skipping the extra operation of adding and removing the item later, which is useful for replicas with no initial state, such as newly added replicas).
- If the item is found in the received rem\_set and also in the local main\_set, it gets removed from the main\_set.
- After iterating over the received add\_set, it iterates over the received rem\_set and removes items from the main\_set, if found.

- Whenever an item is inserted or removed from the `main_set`, the item gets inserted into the local `add_set` or `rem_set` accordingly.

In case of the failure of a node or if a node is newly added, they will eventually catch up with the same final state as other replicas by the above implementation.

The above points prove the convergence property, i.e., eventually, the `main_set` ends up in the same state irrespective of node failure or newly added nodes.

## Commutativity

This property ensures that no matter the order in which the set operations (insertion and deletion of integer values) are performed upon receiving by the replicas, it won't affect the final state. For example, inserting two items into the set in any order won't affect the final state of the set.

### Proof:

- Let's say there are three nodes: A, B, and C.
- A gets a request to insert element 5 into the `main_set`. 5 gets inserted into the `main_set` and also into the `add_set`.
- At the same time, B gets a request to insert element 8 into the `main_set`. 8 gets inserted into the `main_set` and also into the `add_set`.
- Both A and B broadcast their states.
- A's `main_set` now contains {5, 8}
- B's `main_set` now contains {5, 8}
- C's `main_set` also contains {5, 8}, no matter in what order it received the UDP packets from A and B. Eventually, the insertion of the items into its `main_set` will end up in the same final state and won't affect it. The implementation as described ensures this.

## Associativity

It ensures that rearranging or grouping the operations won't affect the final state. In this implementation, it always holds the associativity property.

### Proof:

- Operations on sets are inherently associative.
- For example, applying  $(\text{insert}(a), \text{insert}(b))$  or  $\text{insert}(a), (\text{insert}(b))$  results in the same final state  $\{a, b\}$ .
- Similarly, for removal operations:  $(\text{remove}(a), \text{remove}(b))$  is equivalent to  $\text{remove}(a)$ , then  $\text{remove}(b)$ , resulting in the same state.

## Idempotency

This property ensures that performing the same operation multiple times won't affect the final state.

### Proof:

- If an item is already present in the main\_set, another insertion request for the same item won't affect the state of the set. Upon receiving an insertion request via UDP by the add\_set, it won't affect the final state as the implementation first checks if the element to be inserted via the UDP packet is already present in the local add\_set or not. Whenever an item gets inserted into the main\_set, it also gets inserted into the local add\_set.
- Similarly, for the rem\_set in the received UDP packet, it gets compared with the local rem\_set, and only items not present in the local rem\_set are stepped up for removal from the main\_set.
- Applying the same operation on the set is inherently idempotent. In our case, we ensure idempotency on the basis of CRDT, i.e., upon receiving multiple UDP packets with no changes at all in the add\_set or rem\_set, it isn't going to affect the final state of our sets.

## Handling Reinsertions

- It is to be noted that, as we're using two different sets to maintain the insertion and deletion of items from the main\_set, it won't allow us to insert the same item twice. For example, if an item is added first and then removed from the main\_set, we won't be able to insert the same item again. To solve this issue, we're attaching a unique ID with each item in the set.
- Whenever an item is inserted using the /add endpoint, it gets linked with a unique ID and the pair gets inserted into the main\_set and the add\_set. When removing the item, the rem\_set also stores the unique ID along with the item.
- This removes the conflict, ensuring that an item, irrespective of removal, can be inserted again whenever required or requested.